


Deadlock

Edited slides from <http://cs162.eecs.Berkeley.edu>

Starvation vs Deadlock



- Starvation vs. Deadlock
 - Starvation: thread waits indefinitely
 - » Example, low-priority thread waiting for resources constantly in use by high-priority threads
 - Deadlock: circular waiting for resources
 - » Thread A owns Res 1 and is waiting for Res 2
 - » Thread B owns Res 2 and is waiting for Res 1



- Deadlock \Rightarrow Starvation but not vice versa
 - » Starvation can end (but doesn't have to)
 - » Deadlock can't end without external intervention

Conditions for Deadlock

- Deadlock not always deterministic – Example 2 mutexes:

Thread A

x.P();

y.P();

y.V();

x.V();

Thread B

y.P();


x.P();

x.V();

y.V();

- Deadlock won't always happen with this code
 - » Have to have exactly the right timing (“wrong” timing?)
 - » So you release a piece of software, and you tested it, and there it is, controlling a nuclear power plant...
- Deadlocks occur with multiple resources
 - Means you can't decompose the problem
 - Can't solve deadlock for each resource independently
- Example: System with 2 disk drives and two threads
 - Each thread needs 2 disk drives to function
 - Each thread gets one disk and waits for another one


Bridge Crossing Example




- Each segment of road can be viewed as a resource
 - Car must own the segment under them
 - Must acquire segment that they are moving into
- For bridge: must acquire both halves
 - Traffic only in one direction at a time
 - Problem occurs when two cars in opposite directions on bridge: each acquires one segment and needs next
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
 - Several cars may have to be backed up
- Starvation is possible
 - East-going traffic really fast \Rightarrow no one goes west

Train Example (Wormhole-Routed Network)

- Circular dependency (Deadlock!)
 - Each train wants to turn right
 - Blocked by other trains
 - Similar problem to multiprocessor networks
- Fix? Imagine grid extends in all four directions
 - Force ordering of channels (tracks)
 - » Protocol: Always go east-west first, then north-south
 - Called “dimension ordering” (X then Y)



Dining Lawyers Problem



- Five chopsticks/Five lawyers (really cheap restaurant)
 - Free-for all: Lawyer will grab any one they can
 - Need two chopsticks to eat
- What if all grab at same time?
 - Deadlock!
- How to fix deadlock?
 - Make one of them give up a chopstick (Hah!)
 - Eventually everyone will get chance to eat
- How to prevent deadlock?
 - Never let lawyer take last chopstick if no hungry lawyer has two chopsticks afterwards


Four requirements for Deadlock

- Mutual exclusion
 - Only one thread at a time can use a resource.
- Hold and wait
 - Thread holding at least one resource is waiting to acquire additional resources held by other threads
- No preemption
 - Resources are released only voluntarily by the thread holding the resource, after thread is finished with it
- Circular wait
 - There exists a set $\{T_1, \dots, T_n\}$ of waiting threads
 - » T_1 is waiting for a resource that is held by T_2
 - » T_2 is waiting for a resource that is held by T_3
 - » ...
 - » T_n is waiting for a resource that is held by T_1

Resource-Allocation Graph

- System Model

- A set of Threads T_1, T_2, \dots, T_n
- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has W_i instances
- Each thread utilizes a resource as follows:
 - » Request () / Use () / Release ()




- Resource-Allocation Graph:


- V is partitioned into two types:
 - » $T = \{T_1, T_2, \dots, T_n\}$, the set threads in the system.
 - » $R = \{R_1, R_2, \dots, R_m\}$, the set of resource types in system
- request edge – directed edge $T_i \rightarrow R_j$
- assignment edge – directed edge $R_j \rightarrow T_i$

Resource Allocation Graph Examples


- Recall:
 - request edge – directed edge $T_i \rightarrow R_j$
 - assignment edge – directed edge $R_j \rightarrow T_i$



Simple Resource
Allocation Graph




Allocation Graph
With Deadlock



Allocation Graph
With Cycle, but
No Deadlock


Methods for Handling Deadlocks



- Allow system to enter deadlock and then recover
 - Requires deadlock detection algorithm
 - Some technique for forcibly preempting resources and/or terminating tasks
- Ensure that system will *never* enter a deadlock
 - Need to monitor all lock acquisitions
 - Selectively deny those that *might* lead to deadlock
- Ignore the problem and pretend that deadlocks never occur in the system
 - Used by most operating systems, including UNIX

Deadlock Detection Algorithm

- Only one of each type of resource \Rightarrow look for loops
 - More General Deadlock Detection Algorithm
 - Let $[X]$ represent an m-ary vector of non-negative integers (quantities of resources of each type):
 - $[FreeResources]$: Current free resources each type
 - $[Request_x]$: Current requests from thread X
 - $[Alloc_x]$: Current resources held by thread X
 - See if tasks can eventually terminate on their own
- ```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
 done = true
 Foreach node in UNFINISHED {
 if ([Requestnode] <= [Avail]) {
 remove node from UNFINISHED
 [Avail] = [Avail] + [Allocnode]
 done = false
 }
 }
} until(done)
```
- Nodes left in **UNFINISHED**  $\Rightarrow$  deadlocked



# What to do when detect deadlock?

---

- Terminate thread, force it to give up resources
  - In Bridge example, Godzilla picks up a car, hurls it into the river. Deadlock solved!
  - Shoot a dining lawyer
  - But, not always possible – killing a thread holding a mutex leaves world inconsistent
- Preempt resources without killing off thread
  - Take away resources from thread temporarily
  - Doesn't always fit with semantics of computation
- Roll back actions of deadlocked threads
  - Hit the rewind button on TiVo, pretend last few minutes never happened
  - For bridge example, make one car roll backwards (may require others behind him)
  - Common technique in databases (transactions)
  - Of course, if you restart in exactly the same way, may reenter deadlock once again
- Many operating systems use other options

# Techniques for Preventing Deadlock

---


- Infinite resources [break Mutual exclusion, Hold and wait]
  - Include enough resources so that no one ever runs out of resources. Doesn't have to be infinite, just large  
(e.g. Printer SPOOL [break Mutual exculsion ])
  - Give illusion of infinite resources  
(e.g. virtual memory [break No preemption] )
  - Examples:
    - » Bay bridge with 12,000 lanes. Never wait!
    - » Infinite disk space (not realistic yet?)
- No Sharing of resources (totally independent threads) [break Mutual exclusion]
  - Not very realistic
- Don't allow waiting [break Hold and wait]
  - How the phone company avoids deadlock
    - » Call to your Mom in Toledo, works its way through the phone lines, but if blocked get busy signal.
  - Technique used in Ethernet/some multiprocessor nets
    - » Everyone speaks at once. On collision, back off and retry
  - Inefficient, since have to keep retrying
    - » Consider: driving to San Francisco; when hit traffic jam, suddenly you're transported back home and told to retry!

# Techniques for Preventing Deadlock (cont'd)

- Make all threads request everything they'll need at the beginning. [break Hold and wait]
  - Problem: Predicting future is hard, tend to over-estimate resources
  - Example:
    - » If need 2 chopsticks, request both at same time
    - » Don't leave home until we know no one is using any intersection between here and where you want to go; only one car on the Bay Bridge at a time
- Force all threads to request resources in a particular order preventing any cyclic use of resources [break Circular wait]
  - Thus, preventing deadlock
  - Example (x.P, y.P, z.P,...)
    - » Make tasks request disk, then memory, then...
    - » Keep from deadlock on freeways around SF by requiring everyone to go clockwise

# Review: Train Example (Wormhole-Routed Network)

- Circular dependency (Deadlock!)
  - Each train wants to turn right
  - Blocked by other trains
  - Similar problem to multiprocessor networks
- Fix? Imagine grid extends in all four directions
  - Force ordering of channels (tracks)
    - » Protocol: Always go east-west first, then north-south
  - Called “dimension ordering” (X then Y)



# Avoiding Deadlock

---

- Not by imposing arbitrary rules on processes
- But by carefully analyzing each resource request to see it could be safely granted



# Banker's Algorithm for Avoiding Deadlock

- Toward right idea:
  - State maximum (max) resource needs in advance
  - Allow particular thread to proceed if:  
(available resources - #requested)  $\geq$  max  
remaining that might be needed by any thread
- Banker's algorithm (less conservative):
  - Allocate resources dynamically
    - » Evaluate each request and grant if some  
ordering of threads is still deadlock free afterward
    - » Technique: pretend each request is granted, then run deadlock detection  
algorithm, substituting  
 $([Max_{node}] - [Alloc_{node}] \leq [Avail]) \text{ for } ([Request_{node}] \leq [Avail])$   
Grant request if result is deadlock free (conservative!)



# Banker's Algorithm for Avoiding Deadlock

- ```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```
-



- » Technique: pretend each request is granted, then run deadlock detection algorithm, substituting
 $([Max_{node}] - [Alloc_{node}] \leq [Avail])$ for $([Request_{node}] \leq [Avail])$
Grant request if result is deadlock free (conservative!)

Banker's Algorithm for Avoiding Deadlock

- ```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
 done = true
 Foreach node in UNFINISHED {
 if ($[Max_{node}] - [Alloc_{node}] \leq [Avail]$) {
 remove node from UNFINISHED
 $[Avail] = [Avail] + [Alloc_{node}]$
 done = false
 }
 }
} until(done)
```
- 




- » Technique: pretend each request is granted, then run deadlock detection algorithm, substituting  
 $([Max_{node}] - [Alloc_{node}] \leq [Avail])$  for  $([Request_{node}] \leq [Avail])$   
Grant request if result is deadlock free (conservative!)

# Banker's Algorithm for Avoiding Deadlock

- Toward right idea:
  - State maximum resource needs in advance
  - Allow particular thread to proceed if:  
(available resources - #requested)  $\geq$  max remaining that might be needed by any thread
- Banker's algorithm (less conservative):
  - Allocate resources dynamically
    - » Evaluate each request and grant if some ordering of threads is still deadlock free afterward
    - » Technique: pretend each request is granted, then run deadlock detection algorithm, substituting  
 $([Max_{node}] - [Alloc_{node}] \leq [Avail])$  for  $([Request_{node}] \leq [Avail])$   
Grant request if result is deadlock free (conservative!)
    - » Keeps system in a "SAFE" state, i.e. there exists a sequence  $\{T_1, T_2, \dots, T_n\}$  with  $T_1$  requesting all remaining resources, finishing, then  $T_2$  requesting all remaining resources, etc..
  - Algorithm allows the sum of maximum resource needs of all current threads to be greater than total resources



# Banker's Algorithm Example



- Banker's algorithm with dining lawyers
  - “Safe” (won't cause deadlock) if when try to grab chopstick either:
    - » Not last chopstick
    - » Is last chopstick but someone will have two afterwards
  - What if k-handed lawyers? Don't allow if:
    - » It's the last one, no one would have k
    - » It's 2<sup>nd</sup> to last, and no one would have k-1
    - » It's 3<sup>rd</sup> to last, and no one would have k-2
    - » ...



# Banker's algorithm for a single resource

- Three resource allocation states

| Name    | Used | Maximum |
|---------|------|---------|
| Andy    | 0    | 6       |
| Barbara | 0    | 5       |
| Marvin  | 0    | 4       |
| Suzanne | 0    | 7       |

Available: 10

**Safe**

| Name    | Used | Maximum |
|---------|------|---------|
| Andy    | 1    | 6       |
| Barbara | 1    | 5       |
| Marvin  | 2    | 4       |
| Suzanne | 4    | 7       |

Available: 2

**Safe**

| Name    | Used | Maximum |
|---------|------|---------|
| Andy    | 1    | 6       |
| Barbara | 2    | 5       |
| Marvin  | 2    | 4       |
| Suzanne | 4    | 7       |


Available: 1

**Unsafe**

- State: a list of customers showing the money already loaned and the maximum credit available
- Safe state: if there exists a sequence of other states that leads to all the customers getting loans up to their credit limits

# Resource Trajectories

- Two process resource trajectories
  - At point  $t$  the only safe thing to do is run process A until it gets to  $l_4$



- But...
  - This graphical model is difficult to apply to the general case of an arbitrary number of processes and an arbitrary number of resource classes, each with multiple instances.

# Banker's algorithm for multiple resources

| Process | Tape drives | Plotters | Printers | CD-ROMS |
|---------|-------------|----------|----------|---------|
| A       | 3           | 0        | 1        | 1       |
| B       | 0           | 1        | 0        | 0       |
| C       | 1           | 1        | 1        | 0       |
| D       | 1           | 1        | 0        | 1       |
| E       | 0           | 0        | 0        | 0       |

Resources assigned

| Process | Tape drives | Plotters | Printers | CD-ROMS |
|---------|-------------|----------|----------|---------|
| A       | 1           | 1        | 0        | 0       |
| B       | 0           | 1        | 1        | 2       |
| C       | 3           | 1        | 0        | 0       |
| D       | 0           | 0        | 1        | 0       |
| E       | 2           | 1        | 1        | 0       |

Resources still needed

E = (6342)

P = (5322)

A = (1020)

- E: existing resources
- P: the possessed resources
- A: the available resources
- But ...
  - It is essentially useless because processes rarely know what their maximum resource needs will be in advance.



# Summary

---

- Starvation vs. Deadlock
  - Starvation: thread waits indefinitely
  - Deadlock: circular waiting for resources
- Four conditions for deadlocks
  - Mutual exclusion
    - » Only one thread at a time can use a resource
  - Hold and wait
    - » Thread holding at least one resource is waiting to acquire additional resources held by other threads
  - No preemption
    - » Resources are released only voluntarily by the threads
  - Circular wait
    - »  $\exists$  set  $\{T_1, \dots, T_n\}$  of threads with a cyclic waiting pattern

# Summary (2)

---

- Techniques for addressing Deadlock
  - Allow system to enter deadlock and then recover
  - Ensure that system will never enter a deadlock
  - Ignore the problem and pretend that deadlocks never occur in the system
- Deadlock detection
  - Attempts to assess whether waiting graph can ever make progress
- Deadlock prevention
  - Break one of the four conditions
  - Overly restrictive
- Deadlock avoidance
  - Assess, for each allocation, whether it has the potential to lead to deadlock
  - Banker's algorithm gives one way to assess this
  - Requires information that is usually not available